

mmn_11_computer_vision

November 15, 2018

1 MMN-11

Roman Smirnov

1.1 Preparatory Setup

1.1.1 Library Imports

```
In [2]: import numpy as np
import cv2
import matplotlib.pyplot as plt
%matplotlib inline
```

1.2 Preparatory Setup

1.2.1 PyPlot Defaults Configuration

```
In [3]: plt.rcParams['figure.figsize'] = [14,8]
plt.rcParams['figure.frameon'] = False
plt.rcParams['figure.edgecolor'] = 'none'
plt.rcParams['figure.facecolor'] = 'none'
plt.rcParams['ytick.left'] = False
plt.rcParams['ytick.right'] = False
plt.rcParams['xtick.bottom'] = False
plt.rcParams['xtick.top'] = False
plt.rcParams['ytick.labelleft'] = False
plt.rcParams['ytick.labelright'] = False
plt.rcParams['xtick.labelbottom'] = False
plt.rcParams['xtick.labeltop'] = False
plt.rcParams['axes.grid'] = False
plt.rcParams['image.interpolation'] = 'bilinear'
```

1.3 Question 1

```
In [325]: img = cv2.imread('spock.jpg', cv2.IMREAD_COLOR)[...,:-1] # converted to rgb on the
plt.imshow(img)
```

```
Out[325]: <matplotlib.image.AxesImage at 0x11cc22d30>
```



```
In [326]: # convert to grayscale monochrome since that's the input format for the algorithms b
          gray = cv2.cvtColor(img, cv2.COLOR_RGB2GRAY)
          plt.imshow(gray, cmap='gray')
```

```
Out [326]: <matplotlib.image.AxesImage at 0x11cd195c0>
```



1.3.1 Canny Edge Detector

```
In [327]: edges = cv2.Canny(gray, 89, 144)
          plt.imshow(edges, cmap='gray')
```

```
Out[327]: <matplotlib.image.AxesImage at 0x119579780>
```



1.3.2 Harris Corner Detection

```
In [328]: ##### perform corner detection
corners = cv2.cornerHarris(gray,blockSize=2,ksize=3,k=0.08)


# discard detection values below threshold
# rationale: large positive detection values indicate corners
DETECTION_THRESHOLD = 1.0e-4
corners = np.where(corners>1.0e-4,255,0)

# discard sparse detections, and group together dense detection regions
# rationale: the eye corner areas are apparently dense with various corners
corners = cv2.medianBlur(corners.astype(np.uint8), 3)

# dilate detections for a clearer visualization
corners = cv2.dilate(corners, kernel=np.ones((13,13)), iterations=1)

# visualize detection map
plt.imshow(corners, cmap='gray')
```

Out [328]: <matplotlib.image.AxesImage at 0x11c803320>



```
In [329]: # expand the monochrome image to rgb by duplicating the gray channel
corners_overlay = np.repeat(gray[:, :, None], repeats=3, axis=2)

# mark detection points with blue overlay
corners_overlay[..., 2] = np.maximum(corners_overlay[..., 2], corners)

# visualize grayscale image with blue overlay detection marks
plt.imshow(corners_overlay)
```

Out [329]: <matplotlib.image.AxesImage at 0x11931e048>



Sadly, pointed ear detections are not as easy to isolate :)

1.3.3 Scale Invariant Feature Transform (SIFT)

```
In [330]: # extract corner point coordinates by finding centroids of connected components
eye_corner_coordinates = cv2.connectedComponentsWithStats(corners)[3][1:,:]

# convert coordinates to Keypoint type
eye_corner_keypoints = [cv2.KeyPoint(crd[0], crd[1], 13) for crd in eye_corner_coordinates]

# compute SIFT descriptors from corner keypoints
sift = cv2.xfeatures2d.SIFT_create()
eye_corner_descriptors = [sift.compute(gray, [kp])[1] for kp in eye_corner_keypoints]

print(eye_corner_descriptors)

[array([[ 5.,  83.,  90.,  6.,  7.,  44.,  22.,  2.,  7.,
 113., 113.,  3.,  4., 16.,  5.,  2., 13., 107.,
 51.,  6.,  8.,  2.,  3.,  4., 13., 42., 25.,
 12., 10.,  5.,  6.,  3.,  8., 13., 111., 33.,
  9., 74., 81., 48., 14., 75., 77., 27., 17.,
 75., 113., 61., 22., 113., 95.,  3.,  3., 17.,
```

```

74., 17., 86., 54., 37., 9., 8., 5., 6.,
31., 9., 4., 4., 3., 1., 33., 113., 44.,
9., 5., 12., 19., 15., 54., 113., 36., 43.,
6., 5., 19., 40., 81., 59., 30., 113., 4.,
0., 2., 18., 7., 4., 76., 4., 4., 3.,
4., 3., 9., 14., 6., 1., 3., 10., 62.,
24., 7., 4., 3., 66., 5., 2., 74., 40.,
12., 1., 14., 113., 8., 0., 2., 13., 2.,
0., 28.]], dtype=float32), array([[ 15., 26., 29., 22., 8., 4., 4
14., 97., 44., 4., 1., 4., 4., 4., 7.,
112., 85., 3., 2., 5., 6., 5., 4., 112.,
57., 7., 3., 13., 28., 54., 7., 8., 23.,
74., 24., 13., 32., 1., 3., 112., 96., 70.,
57., 97., 39., 10., 38., 112., 76., 10., 112.,
112., 52., 19., 95., 101., 15., 3., 33., 88.,
72., 50., 2., 0., 3., 64., 18., 10., 31.,
10., 3., 1., 3., 35., 70., 102., 35., 5.,
13., 17., 8., 9., 98., 112., 33., 11., 27.,
22., 8., 8., 75., 112., 30., 37., 5., 0.,
3., 47., 6., 1., 8., 20., 16., 1., 1.,
21., 20., 5., 4., 9., 4., 2., 1., 5.,
16., 21., 9., 3., 2., 2., 3., 3., 7.,
21., 9.]], dtype=float32)]

```

1.3.4 SIFT Pattern Matching

In [331]: *# load and visualize images*

```

img1 = cv2.imread('spock.jpg', cv2.IMREAD_GRAYSCALE)
img2 = cv2.imread('spock2.jpg', cv2.IMREAD_GRAYSCALE)
ax1, ax2 = plt.subplots(1,2)[1]
ax1.imshow(img1, cmap='gray'); ax2.imshow(img2, cmap='gray')

```

Out[331]: <matplotlib.image.AxesImage at 0x11d58e9e8>



```


In [332]: # detect keypoints and compute descriptors
          kp1, des1 = sift.detectAndCompute(img1, None)
          kp2, des2 = sift.detectAndCompute(img2, None)

In [333]: # compute descriptor pattern matches
          brute_force_matcher = cv2.BFMatcher(crossCheck=True)
          matches = brute_force_matcher.match(des1, des2)
          # sort matches by way of matcher default distance function (l2 norm)
          matches = sorted(matches, key = lambda x:x.distance)

In [335]: # visualize top 10 matches
          matches_img = cv2.drawMatches(img1,kp1,img2,kp2,matches[:10],None, flags=2)
          plt.imshow(matches_img)

Out[335]: <matplotlib.image.AxesImage at 0x11cbe2c18>

```



1.3.5 Hough Line Transform

```

In [15]: # load as BGR, convert to RGB, convert to grayscale, showcase
          phasers = cv2.imread('enterprise_phasers.jpg', cv2.IMREAD_COLOR)[...::-1]
          phasers_gray = cv2.cvtColor(phasers, cv2.COLOR_RGB2GRAY)
          ax1, ax2 = plt.subplots(2,1)[1]
          ax1.imshow(phasers); ax2.imshow(phasers_gray, cmap='gray')


Out[15]: <matplotlib.image.AxesImage at 0x1133628d0>

```




```
In [16]: # isolate phaser rays
         phaser_blur = cv2.blur(phasers_gray, (7,7))
         phaser_edges = cv2.Canny(phaser_blur,128,256)
         phaser_edges = cv2.morphologyEx(phaser_edges, cv2.MORPH_CLOSE, np.ones((7,7)))
         plt.imshow(phaser_edges, cmap='gray')
```

```
Out[16]: <matplotlib.image.AxesImage at 0x112e14588>
```




```
In [17]: # find diagonal line terminal point coordinates
lines = cv2.HoughLinesP(phaser_edges, rho=1, theta=1.5*np.pi/180, threshold=256,
                        minLineLength=128, maxLineGap=64)

# draw found lines on image
phaser_lines = np.copy(phasers)
for line in lines:
    x1,y1,x2,y2 = line[0]
    cv2.line(phaser_lines,(x1,y1),(x2,y2),(0,255,0),7)

# visualize result
plt.imshow(phaser_lines)
```

```
Out[17]: <matplotlib.image.AxesImage at 0x112dc1a20>
```



1.4 Question 2

1.4.1 K-Means Clustering

We'll cluster 8 random points (without replacement) from the range [1,12] into 3 different groups.

In [18]: *# prints some data, plots cluster assignments*

```
def plot_kmeans_step(pts, lbls, i, cntrs):
    plt.figure()
    plot = np.zeros((13,13,3),np.uint8)
    colors = np.where(lbls==0, (255,0,0), lbls)
    colors = np.where(lbls==1, (0,255,0), colors)
    colors = np.where(lbls==2, (0,0,255), colors)
    pts_y, pts_x = pts.astype(np.uint8).T
    plot[pts_y, pts_x] = colors
    plt.imshow(plot)
    plt.suptitle('iteration='+str(i)+' , centers=\n'+str(cntrs))
```

In [17]: *# select random 2d coordinates from valid range*

```
pts = np.random.permutation(np.mgrid[1:13, 1:13].reshape(2, -1).T)[:8].astype(np.float)
# randomly set initial kmeans cluster assignments
lbls = np.random.choice(3,(8,1)).astype(np.int32)
# plot initial assignment
plot_kmeans_step(pts, lbls, 0, 'not computed yet')


# previous iteration's cluster centers
old_cntrs=None
```

```


for i in range(100):
    # run kmeans for 1 step
    ret, lbls, cntrs = cv2.kmeans(pts,
                                   K=3,
                                   bestLabels=lbls,
                                   criteria=(cv2.TERM_CRITERIA_MAX_ITER, 1, 1.0),
                                   attempts=1,
                                   flags=cv2.KMEANS_USE_INITIAL_LABELS)
    # print some data, plot current cluster assignments
    plot_kmeans_step(pts, lbls, i, cntrs)
    # terminate loop when centers didn't change from last iteration
    if np.array_equal(cntrs, old_cntrs):
        break;
    old_cntrs=cntrs

```


iteration=0, centers=not computed yet



```
iteration=0, centers=[[ 5.33333349 10.66666698]
 [ 3.      6.33333349]
 [11.5     6.      ]]
```



```
iteration=1, centers=[[ 5.33333349 10.66666698]
[ 3.      6.33333349]
[ 11.5    6.      ]]
```



1.5 Question 3

1.5.1 Least Squares Polynomial Fitting

Overview We'll fit a parabolic curve to a set of 2d points using the least squares method. The least squares method works by finding a curve such that the sum of the squared errors between the given coordinates and the curve is minimal.

Math The error is defined as

$$E = \sum_{i=0}^n (f(x_i) - y_i)^2$$

Where $f(x)$ is our parabola fitting function and x_i, y_i are the input coordinates.
We'll minimize the error by computing the vandermonde matrix

$$\mathbf{y} = \mathbf{X}\mathbf{a}$$

Where \mathbf{y} are the y coordinate column vector, \mathbf{X} is the x coordinate polynomial power matrix, and \mathbf{a} is the sought coefficient column vector.

Note: see [wolfram least squares polynomial fitting](#)

```
In [299]: # the specificed input coordinates
pts = np.array(((1,3.96), (4,27.96), (3,15.15), (5,45.8), (2,7.07), (6,69.4)))
pts

Out[299]: array([[ 1. ,  3.96],
 [ 4. , 27.96],
 [ 3. , 15.15],
 [ 5. , 45.8 ],
 [ 2. ,  7.07],
 [ 6. , 69.4 ]])

In [109]: # calculate the x coordinate powers matrix
pts_x_pow = np.stack(((pts[:,0]**2).T, pts[:,0].T, np.ones(pts.shape[0]).T)).T
# solve the linear equation
dst = cv2.solve(pts[:,1],x, flags=cv2.DECOMP_NORMAL)[1]
```

Note: see [decomposition types](#)

```
In [110]: # set curve x values
curve_x = np.arange(13)
# calculate curve y values
curve_x_pow = np.stack(((curve_x**2).T, curve_x.T, np.ones(curve_x.shape[0]).T)).T
curve_y = np.matmul(curve_x_pow,dst.T).ravel().astype(np.uint8)

In [111]: # plot input points and points on calculated curve
pts_x, pts_y = pts.astype(np.uint8).T
plot = np.zeros((89,21), np.uint8)
plot[pts_y, pts_x] = 255; plot[curve_y, curve_x] = 100
plt.imshow(plot[::-1,:], cmap='gray')

Out[111]: <matplotlib.image.AxesImage at 0x1187fd0>
```




1.6 Question 4

1.6.1 Principal Component Analysis (PCA)

```
In [300]: # set input points
pts = np.array(((2.5,2.9),(0.5,1.2),(2.2,3.4),(1.9,2.7),(3.1,3.5),(2.3,3.2),(2,2.1),
                (1,1.6),(1.5,2.1),(1.1,1.4))).T


pts.T

Out[300]: array([[ 2.5,  2.9],
                 [ 0.5,  1.2],
```

```
[ 2.2,  3.4],  
[ 1.9,  2.7],  
[ 3.1,  3.5],  
[ 2.3,  3.2],  
[ 2. ,  2.1],  
[ 1. ,  1.6],  
[ 1.5,  2.1],  
[ 1.1,  1.4]])
```

```
In [302]: # plot input points  
plot = np.zeros((8,8), np.uint8)  
pts_x, pts_y = pts.astype(np.uint8)  
plot[pts_y, pts_x] = 255  
plt.imshow(plot[::-1,:], cmap='gray', interpolation='none')
```

```
Out[302]: <matplotlib.image.AxesImage at 0x11b98d5f8>
```



```
In [303]: # calculate covariance matrix, its eigenvalues and eigenvectors
          eigval, eigvec = np.linalg.eig(np.cov(pts))
          print(' eigenvalues: ', eigval)
          print(' eigenvectors: \n', eigvec) # cols

eigenvalues: [ 0.0490834  1.28402771]
eigenvectors:
[[-0.73517866 -0.6778734 ]
 [ 0.6778734  -0.73517866]]
```


Note: the principal components are defined as those having the largest eigenvalues.

```
In [304]: # calculate projection into a 1d line
          proj = np.matmul(eigvec.T[1], pts)*-1 # negative relation
          proj

Out[304]: array([ 3.8267016 ,  1.22115109,  3.99092891,  3.27294183,  4.67453283,
                  3.91168051,  2.89962197,  1.85415925,  2.56068527,  1.77491086])

In [319]: # plot projected line
          plot = np.zeros((8,1), np.uint8)
          plot[proj.astype(np.uint8), 0] = 150
          plt.imshow(plot[::-1,:], cmap='gray', interpolation='none')

Out[319]: <matplotlib.image.AxesImage at 0x11c2d39b0>
```



```
In [320]: # calculate inverse projection back into 2d coordinates
          inv_proj = np.matmul(eigvec,np.vstack((np.zeros(proj.shape[0]), proj)))*-1
          inv_proj

Out[320]: array([[ 2.59401922,  0.82778584,  2.70534454,  2.2186402 ,  3.16874146,
                   2.65162416,  1.9655766 ,  1.25688523,  1.73582043,  1.20316485],
                  [ 2.81330934,  0.89776421,  2.93404575,  2.40619697,  3.43661676,
                   2.87578402,  2.13174018,  1.3631383 ,  1.88256116,  1.30487658]])

In [321]: # plot backwards transformed result
          plot = np.zeros((8,8), np.uint8)
```

```
plot[inv_proj[1].astype(np.uint8), inv_proj[0].astype(np.uint8)] = 255  
plt.imshow(plot[::-1,:], cmap='gray', interpolation='none')
```

Out[321]: <matplotlib.image.AxesImage at 0x11c303278>

